

FH Aachen

Department of Electrical Engineering and Information Technology

Exposé

**Evaluating AI Agent Capabilities in SysML v2 Modeling
via an MCP Bridge to Magic Systems of Systems Architect**

Nassim Awabdy

nassim.awabdy@alumni.fh-aachen.de

Supervisor: Prof. Dr. rer. nat. Sebastian Voss

Second examiner: M. Eng. Tim Höner

Master's Thesis – WS 2026/2027
Aachen, September 2026

1 Exposé

Motivation

[TODO 150 words]

Problem Statement

[TODO 150 words]

Research Question and Hypotheses

RQ: Which SysML v2 modeling tasks can an LLM-based agent perform reliably in an industrial modeling environment (MSoSA) when connected through an MCP bridge, and where do its capabilities break down?

- **H1 (Task asymmetry):** Read-only tasks (Abfragen, Validierung) are performed with high reliability, whereas write tasks (Korrektur, Erstellen) and Verifizierung degrade with model size and cross-layer dependency depth.
- **H2 (Tool semantics):** Exposing tool-native capabilities (validation suites, expression evaluation) as MCP tools yields larger gains than a thin CRUD wrapper over the SysML v2 API alone.
- **H3 (Model scale):** Findings obtained on small teaching models (GfSE collection) do not transfer to a large multi-layer model (Apollo 11) without a semantic navigation layer.

Approach

Use case	Example task on Apollo 11	Ground truth / metric
Abfragen (query)	“Which technical components satisfy requirement X?”	exact set match vs. model query
Validierung (validation)	“Is package P well-formed and CoSMA-conformant?”	tool validation suite results
Verifizierung (verification)	“Is mass budget requirement R satisfied by the current values?”	expression evaluation in tool
Korrektur (correction)	repair injected faults (dangling ref., wrong type, broken <code>satisfy</code>)	fault-repair rate, no collateral change
Erstellen (creation)	add missing technical component with ports + satisfy links	syntactic validity + expert rubric

Table 1: Use-case catalogue (to be refined).

[TODO 150 words prose around the table]

Evaluation Methodology

[TODO 150 words]

Bibliography

- [1] OBJECT MANAGEMENT GROUP: OMG Systems Modeling Language (SysML), Version 2.0 / Object Management Group. Version: September 2025. <https://www.omg.org/spec/SysML/2.0/>. 2025. – Specification. – The formal SysML v2 language specification (textual notation, KerML-based semantics). Defines the language the agent must read and write.
- [2] OBJECT MANAGEMENT GROUP: Systems Modeling API and Services, Version 1.0 / Object Management Group. Version: September 2025. <https://www.omg.org/spec/SystemsModelingAPI/1.0/>. 2025. – Specification. – Standard REST/HTTP API for SysML v2 repositories (projects, commits, branches, elements, queries). Implemented by Teamwork Cloud; the transport used by Bazzal et al. and the open-source MCP servers, and the baseline the thesis contrasts with a tool-internal bridge.
- [3] ANTHROPIC: *Model Context Protocol – Specification*. <https://modelcontextprotocol.io/specification>, 2025. – Open protocol for exposing tools, resources and prompts to LLM agents. The bridge to MSoSA is implemented as an MCP server.
- [4] HELLE, Philipp ; SCHRAMM, Gerrit: Fly me to the Moon – Modeling Apollo 11 Using SysML v2. In: *Systems Engineering* (2026), Juni. <http://dx.doi.org/10.1002/sys.70074>. – DOI 10.1002/sys.70074. – [TOP 3] Presents the open-source Apollo 11 SysML v2 model (Airbus Central R&T, MPL-2.0, ≈ 7 kLOC, five CoSMA layers) explicitly intended as a benchmark for next-generation MBSE tools. Method: modelling case study with documented layering, traceability and deliberate gaps. Chosen as the system under test for all use-case experiments; the documented gaps seed the *Validierung* and *Korrektur* tasks.
- [5] GFSE – GESELLSCHAFT FÜR SYSTEMS ENGINEERING: *SysML-v2-Models: A curated collection of high-quality SysML v2 models*. <https://github.com/GfSE/SysML-v2-Models>, accessed 12.09.2026, 2025. – Curated BSD-licensed collection of small SysML v2 models, explicitly meant to help train and test LLM-based tooling. Used as secondary/generalisation test set and as source of good/bad examples.
- [6] DASSAULT SYSTÈMES: *CATIA Magic/Cameo SysML v2 Solution – Documentation*. <https://docs.nomagic.com/SYSML2P/latest/>, accessed 12.09.2026, 2026. – Vendor documentation for SysML v2 in Magic Systems of Systems Architect / Cameo: metamodel conformance, validation suites, expression evaluation, simulation, and the standard SysML v2 REST API served by Teamwork Cloud. Defines what tool-native capabilities the semantic bridge tier can expose; vendor claims are cited as claims.

- [7] AJHCS: *cameo-mcp-bridge: MCP server bridging AI assistants to CATIA Magic / Cameo*. <https://github.com/ajhcs/cameo-mcp-bridge>, accessed 12.09.2026, 2026. – Community MCP server for Cameo (SysML v1/UML) via an embedded Java OpenAPI plugin with 162 tools. Closest engineering precedent for the tool-internal bridge (Option B); no SysML v2 support and no systematic evaluation.
- [8] DEHART, John K.: Leveraging Large Language Models for Direct Interaction with SysML v2. In: *INCOSE International Symposium 34* (2024), Juli, No. 1, p. 2168–2185. <http://dx.doi.org/10.1002/iis2.13262>. – DOI 10.1002/iis2.13262. – The founding paper of the topic: argues that SysML v2’s English-like textual syntax and the maturation of LLMs make conversational interaction with system models feasible, and that LLMs can replace API expertise as the interface. Method: conceptual analysis with case studies; explicitly warns of automation bias. Provides the vision the thesis puts to an empirical test.
- [9] FRESEMAN, Carina ; ELLSEL, Claudius ; GRUNENWALD, Carl-Philipp ; STARK, Rainer: Using Large Language Models for System Modeling: Current State and Future Implications. In: *2025 IEEE International Symposium on Systems Engineering (ISSE)*, IEEE, Oktober 2025, p. 1–8. – Structured review of 20 publications on LLMs for MBSE tasks (creation, analysis, modification, reformulation), plus a modular prompting framework for SysML v1→v2 translation. Points to MCP as the emerging mechanism for supplying engineering context. Provides the task taxonomy on which the use-case catalogue (Abfragen, Validierung, ...) is mapped.
- [10] DEHN, Simon ; SCHNÜRER, Simon ; JACOBS, Georg ; HÖPFNER, Gregor: Generating SysML v2 Models from Natural Language Requirements Using Large Language Models. In: *2025 IEEE International Symposium on Systems Engineering (ISSE)*, IEEE, Oktober 2025, p. 1–7. – Four-component prompting (role, ontology, few-shot, requirements) for NL→SysML v2, evaluated on an automotive electric-window case with precision/recall/F1 and syntax metrics. Reference point for the *Erstellen* use case and its element-level metrics; file-based generation without a modelling tool.
- [11] QUALIS, Richard A.: Mitigating Hallucinations in SysML v2 Generation Using LLMs and a Tri-Layered Knowledge Graph Reasoning Framework. In: *2025 ACM/IEEE 28th International Conference on Model Driven Engineering Languages and Systems Companion (MODELS-C)*, IEEE, Oktober 2025, p. 357–366. – Three-layer knowledge graph (SysML patterns / domain constructs / system-specific) to ground SysML v2 generation and reduce hallucinated constructs; satellite case study measuring structural validity and semantic fidelity. Motivates grounding the agent in the *live* model instead of a hand-built static KG.

- [12] WU, Rui ; WEN, Guangwei ; ZHANG, XiaoNing ; LIANG, Xiaojun ; YANG, Chunhua ; GUI, Weihua: Knowledge-Augmented Conversational Agents for Dependency-Aware SysML v2 Model Generation and Refinement. In: *2025 IEEE 26th China Conference on System Simulation Technology and its Applications (CCSSTA)*, IEEE, Juli 2025, p. 581–586. – SysForge: multi-agent framework (conversational, context-synthesizer, designer and validator agents) with dependency-aware KG retrieval and iterative refinement; beats plain LLM and semantic-RAG baselines on Pass@1 and BLEU. Closest prior work for *Korrektur* through a generate–validate loop; file-based and without an industrial tool.
- [13] QUAST, Vincent ; JACOBS, Georg ; DEHN, Simon ; HÖPFNER, Gregor: Enabling Humans and AI Systems to Retrieve Information from System Architectures in Model-Based Systems Engineering. In: *Systems* 14 (2026), Januar, No. 1, p. 83. <http://dx.doi.org/10.3390/systems14010083>. – DOI 10.3390/systems14010083. – GraphRAG multi-agent system (supervisor + graph-query agent) over a Neo4j graph parsed from SysML v2 text and schema'd by RFLP; 93% retrieval accuracy on a battery-EV model. Reference for the *Abfragen* use case and its accuracy metric; the thesis queries the live tool instead of a derived graph.
- [14] BOUAMRA, Yasmine ; YUN, Bruno ; POISSON, Alexandre ; ARMETTA, Frédéric: *SysTemp: A Multi-Agent System for Template-Based Generation of SysML v2*. arXiv:2506.21608 [cs.AI]. – Multi-agent, template-driven generation of SysML v2 from natural-language specifications, motivated by the scarcity of training corpora and the complex syntax. Its evaluation section states openly that no semantic evaluation method and no benchmark exist for SysML v2 generation – the gap citation for the thesis's benchmark contribution.
- [15] AL-SHAMI, Haitham ; MALIK, Rohail ; ALA-LAURINAHO, Riku ; VEPSÄLÄINEN, Jari ; VIITALA, Raine: Automated Semantic Fault Localization in SysML v2: A Human-in-the-Loop Framework Using Knowledge-Graph Augmented LLMs. In: *INCOSE International Symposium* 36 (2026), Juli, No. 1, p. 1274–1286. <http://dx.doi.org/10.1002/iis2.70189>. – DOI 10.1002/iis2.70189. – Fine-tuned small language model plus a domain KG of physical compatibility rules that localises and repairs *semantic* faults which pass SysML v2 syntax validation; synthetic faults injected via the KG, output as unified-diff patches, 1 184 test samples, repair rate <3% → >91%. Direct methodological precedent for the fault-injection design of the *Korrektur* use case and for the syntactic/semantic fault distinction.
- [16] BAZZAL, Mahmoud ; SCHLEIFER, Simon ; LUNGU, Adriana ; KRUSE, Benjamin ; GOETZ, Stefan ; KOLDEWEY, Christian ; TISSEN, Denis ; WARTZACK, Sandro J. ; DUMITRESCU, Roman: Model Context Protocol for MBSE: A Unified Interoperable

- Agentic Framework for Systems Modelling. In: *Proceedings of the Design Society 6* (2026), Juli, p. 2721–2730. <http://dx.doi.org/10.1017/pds.2026.10630>. – DOI 10.1017/pds.2026.10630. – [TOP 3] Fraunhofer IEM / FAU framework for implementing MCP servers on top of the standard SysML v2 API, so that LLMs can operate on system models interoperably across tools. Closest scientific competing work: it delivers the *artifact* class the thesis builds, but (per abstract) without a capability evaluation and via the standard REST API rather than a tool-internal interface. The thesis is positioned as the missing empirical evaluation. Full text to be checked (open access, CC BY-NC-ND).
- [17] AHLBRECHT, Alexander ; ZAESKE, Wanja Marlo M. ; DURAK, Umut: MBSqlE: Enabling SQL as Powerful Query Language for SysML v2 in Aviation. In: *44th AIAA DATC/IEEE Digital Avionics Systems Conference (DASC 2025)*. Montreal, Canada : IEEE, September 2025. – DLR Institute of Flight Systems: shows that the standard SysML v2 API query model offers no direct graph traversal and hence hits performance bottlenecks on large models, and proposes SQLite as an intermediate, traversable representation (open-source `DLR-FT/sysml-v2-sql`). Technical motivation for the semantic/cache layer of the bridge and for querying inside the tool rather than through the REST API.
- [18] LIPARULO, Leonardo ; PIERRI, Francesco: *Benchmarking AI Agents for Hardware Design Automation via MCP Tool Calling*. arXiv:2608.26199 [cs.AI]. – Methodological sibling in an adjacent domain: an MCP server that reproduces the state and dependency logic of a proprietary hardware-design tool, a benchmark of single-operation edits, dependency chains, invalid and misspelled requests, and seven models under pipeline ablations. Findings – comprehensive tool descriptions reduce failures, few-shot can cause inaction, cumulative context hurts constrained models – directly inform the bridge and benchmark design.
- [19] SHI, Jiasheng ; ZHANG, Tianhan: *What Do CAE Simulation Agents Really Need Beyond a Generic Harness?* arXiv:2609.03718 [cs.CE]. – With information access and repair budget held fixed, a single-agent generic harness matches or beats specialised multi-agent CAE systems (FoamBench 96.4 % vs. 88.2 %); execution-feedback repair and domain tutorials explain the gains, scripted reflection adds nothing. Counter-evidence to H2 (semantic bridge > thin bridge) that the ablation must engage.
- [20] WANG, Zizhe: *Beyond Executable Models: The Pufibara Agent Harness and the Modelica Agent Workflow Benchmark for Physical System Modeling*. arXiv:2608.23653 [cs.SE]. – 232-task benchmark (repair, generation, tuning) for Modelica agents, with a source-grounded task-construction method and a benchmark-owned evaluator outside the agent loop; compares two complete harnesses under matched LLM backends.

The external-evaluator principle and the repair/generation/tuning split are adopted for the SysML v2 benchmark.

- [21] SHEFA, Jannatul ; SALADO, Alejandro ; WACH, Paul ; TOPCU, Taylan G.: *Two Truths and A Lie? Benchmarking Off-the-Shelf LLMs for Requirements Quality Assessment: Performance, False Alarms, and Misses*. arXiv:2609.03230 [cs.SE]. – Ten OpenAI/Anthropic models, 100 runs, expert ground truth on INCOSE quality criteria: the best model finds a median of only 47 % of expert-identified issues, judgement-heavy criteria are almost always missed, generational progress is non-monotonic. Justifies expert-defined acceptance criteria and a second human rater instead of LLM-as-judge in the evaluation design.
- [22] PRADAS-GOMEZ, Alejandro ; BRAHMA, Arindam ; ISAKSSON, Ola: *DUCTILE: Agentic LLM Orchestration of Engineering Analysis in Product Development Practice*. arXiv:2603.10249. – Separates adaptive orchestration by the LLM agent from deterministic execution by verified engineering tools, demonstrated on an industrial aerospace structural-analysis task; evaluated against expert-defined acceptance criteria over 10 repeated independent runs and in deployment with engineers. Template for the agent-tool division of labour and for repeated-run reliability reporting.
- [23] MOLINARI, Gioele ; FELTEN, Florian ; MASSOUDI, Soheyl ; FUGE, Mark: *EngiAI: Capability-Based Evaluation of Tool-Connected LLM Agents for Engineering Design*. arXiv:2605.19743 (extended version of an ASME IDETC-CIE 2026 paper). – Capability-based evaluation framework for tool-connected engineering agents that scores workflow execution, retrieval-assisted parameter selection, HPC orchestration and code authoring separately from execution traces and artifacts; four LLM backends on EngiBench problems. Argues for capability-specific scores over end-to-end demos – the rationale for evaluating per use case rather than per end-to-end task.
- [24] HASAN, Mohammed M. ; LI, Hao ; RAJBAHADUR, Gopi K. ; ADAMS, Bram ; HASSAN, Ahmed E.: *Model Context Protocol (MCP) Tool Descriptions Are Smelly! Towards Improving AI Agent Efficiency with Augmented MCP Tool Descriptions*. arXiv:2602.14878. – Empirical study of 856 tools across 103 MCP servers: 97 % of tool descriptions contain at least one smell; augmenting descriptions raises task success by a median 5.85 pp but increases execution steps by 67 %. Supplies the description rubric for the bridge's tool surface and a confound to control between the thin and semantic arms.
- [25] SOFTWARE ENGINEERING INSTITUTE, CARNEGIE MELLON UNIVERSITY: *Native AI Integration for Model-Based Systems Engineering: Three Layers that Make It Work*. SEI Blog, <https://www.sei.cmu.edu/blog/native-ai-integration-for-model-based-systems-engineering-three-layers->

accessed 12.09.2026, 2026. – [TOP 3] Controlled three-arm SysML v2 experiment (no tools / validation loop / MCP knowledge server + skills; 8 tasks, Claude Opus 4.6): iterative validation eliminates syntax errors, curated knowledge raises pattern conformance from 71.7 to 94.1. Grey literature, but the only controlled tooling ablation on SysML v2 to date; file-based, so the thesis differentiates on the live-model / industrial-tool axis and adopts its ablation template.