

FH Aachen

Department of Electrical Engineering and Information Technology

Exposé

**Evaluating AI Agent Capabilities in SysML v2 Modeling
via an MCP Bridge to Magic Systems of Systems Architect**

Nassim Awabdy

nassim.awabdy@alumni.fh-aachen.de

Supervisor: Prof. Dr. rer. nat. Sebastian Voss

Second examiner: M. Eng. Tim Höner

Master's Thesis – WS 2026/2027
Aachen, September 2026

1 Exposé

Motivation

[TODO 150 words]

Problem Statement

[TODO 150 words]

Research Question and Hypotheses

RQ: Which SysML v2 modeling tasks can an LLM-based agent perform reliably in an industrial modeling environment (MSoSA) when connected through an MCP bridge, and where do its capabilities break down?

- **H1 (Task asymmetry):** Read-only tasks (Abfragen, Validierung) are performed with high reliability, whereas write tasks (Korrektur, Erstellen) and Verifizierung degrade with model size and cross-layer dependency depth.
- **H2 (Tool semantics):** Exposing tool-native capabilities (validation suites, expression evaluation) as MCP tools yields larger gains than a thin CRUD wrapper over the SysML v2 API alone.
- **H3 (Model scale):** Findings obtained on small teaching models (GfSE collection) do not transfer to a large multi-layer model (Apollo 11) without a semantic navigation layer.

Approach

Use case	Example task on Apollo 11	Ground truth / metric
Abfragen (query)	“Which technical components satisfy requirement X?”	exact set match vs. model query
Validierung (validation)	“Is package P well-formed and CoSMA-conformant?”	tool validation suite results
Verifizierung (verification)	“Is mass budget requirement R satisfied by the current values?”	expression evaluation in tool
Korrektur (correction)	repair injected faults (dangling ref., wrong type, broken <code>satisfy</code>)	fault-repair rate, no collateral change
Erstellen (creation)	add missing technical component with ports + satisfy links	syntactic validity + expert rubric

Table 1: Use-case catalogue (to be refined).

[TODO 150 words prose around the table]

Evaluation Methodology

[TODO 150 words]

Bibliography

- [1] INTERNATIONAL COUNCIL ON SYSTEMS ENGINEERING (INCOSE): Systems Engineering Vision 2020 / INCOSE. Version: September 2007. https://sdincose.org/wp-content/uploads/2011/12/SEVision2020_20071003_v2_03.pdf. 2007 (INCOSE-TP-2004-004-02, v2.03). – Technical Product. – [TOP 3] The original, canonical coinage of MBSE: "the formalized application of modeling to support system requirements, design, analysis, verification and validation activities beginning in the conceptual design phase and continuing throughout development and later life cycle phases." Nearly every downstream MBSE paper traces its definition to this document; verified by reading the primary PDF directly, not a secondary citation. PDF at <bib/papers/incose2007-se-vision-2020.pdf>.
- [2] OBJECT MANAGEMENT GROUP: Kernel Modeling Language (KerML), Version 1.0 / Object Management Group. Version: März 2026. <https://www.omg.org/spec/KerML/1.0/PDF>. 2026 (formal/26-03-01). – Specification. – Formal specification of KerML, the first-order-logic kernel that SysML v2 is built on (Root/Core/Kernel layers). Method: normative language specification. Foundational reference for why SysML v2 guidelines and constraints can be verified natively as CSPs rather than transformed to external formalisms as in SysML v1; PDF at <bib/papers/omg-kerml-v1.0.pdf>. Note: KerML Version 1.0 was *adopted* September 2025; the cited PDF is OMG document `formal/26-03-01`, published March 2026 (this is the document number/date, not the version-adoption date) – verified against the live spec page <https://www.omg.org/spec/KerML/1.0> on 20.09.2026.
- [3] BOELSEN, Kathrin ; MAY, Matthias ; JACOBS, Georg ; MENNICKEN, Manuel ; MOERS, Frederik ; ZERWAS, Thilo ; HÖPFNER, Gregor: SysML v2 based modelling guidelines for mechanical system elements. In: *Forschung im Ingenieurwesen* 89 (2025), No. 1, p. 60. <http://dx.doi.org/10.1007/s10010-025-00827-w>. – DOI 10.1007/s10010-025-00827-w. – Enforces domain-specific modelling guidelines intrinsically in SysML v2 via Metadata Definitions and the Definition/Usage separation (moteogo methodology, mechanical system elements), instead of extrinsic UML profiles/OCL as in v1. Method: metamodel design plus a worked example (lubricated rolling-contact component). Concrete precedent for how a semantic MCP bridge could expose or enforce guideline conformance as a tool-native check; PDF at <bib/papers/boelsen2025-sysml2-guidelines.pdf>.
- [4] AHLBRECHT, Alexander ; ZAESKE, Wanja ; DURAK, Umut: MBSqLE: Enabling SQL as Powerful Query Language for SysML v2 in Aviation. In: *44th AIAA DATC/IEEE Digital Avionics Systems Conference (DASC 2025)*. Montreal, Canada : IEEE, September 2025. – DLR Institute of Flight Systems: analyses the standard SysML v2 API query

model and shows it offers no direct graph traversal, leading to performance bottlenecks on large models; proposes SQLite as an intermediate, traversable representation with SQL as query language (open-source `DLR-FT/sysml-v2-sql`). Method: API analysis plus design and open-source implementation. Technical motivation for the semantic/cache layer of the bridge and for querying inside the tool rather than through the REST API; PDF at `bib/papers/ahlbrecht2025mbsqle.pdf`.

- [5] FRESEMANN, Carina ; ELLSEL, Claudius ; GRUNENWALD, Carl-Philipp ; STARK, Rainer: Using Large Language Models for System Modeling: Current State and Future Implications. In: *2025 IEEE International Symposium on Systems Engineering (ISSE)*, IEEE, Oktober 2025, p. 1–8. – Structured literature review of 20 publications on LLMs for MBSE tasks (creation, analysis, modification, reformulation), plus a modular prompting framework for SysML v1→v2 translation. Points to MCP as the emerging mechanism for supplying engineering context, potentially surpassing prompting-only methods. Flags the lack of standardized, comparable evaluation criteria as an open problem. Provides the task taxonomy on which the thesis’s use-case catalogue (Abfragen, Validierung, Verifizierung, Korrektur, Erstellen) is mapped. PDF at `bib/papers/fresemann2025review.pdf`.
- [6] DEHART, John K.: Leveraging Large Language Models for Direct Interaction with SysML v2. In: *INCOSE International Symposium 34* (2024), Juli, No. 1, p. 2168–2185. <http://dx.doi.org/10.1002/iis2.13262>. – DOI 10.1002/iis2.13262. – The founding paper of the topic: argues that SysML v2’s English-like textual syntax and the maturation of LLMs make conversational interaction with system models feasible, and that LLMs can replace API expertise as the interface. Method: conceptual analysis plus three OpenAI Assistants-API case studies – (1) modifying a beam model’s length attribute, where the LLM generates and locally executes ad-hoc Python/regex code against the model’s text file (no structured tool/API call); (2) a Flask `/ask` endpoint wrapping the same Assistant for conversational querying, e.g. correctly computing a von Mises stress from boundary conditions given only in prose; (3) a Jupyter-notebook case validating two requirements (max volume, max length) against a block model with a deliberately unconventional calculation, correctly traversing the Definition/Usage relationship and flagging the volume requirement as failed. All three are single-run proofs of concept – no ground truth, no repeated trials, no quantitative metric beyond the reported outcome – and explicitly warns of automation bias and overreliance. Origin of the idea the thesis puts to an empirical test; also the earliest documented instance of the “no structured tool, LLM edits text directly” pattern that motivates this thesis’s no-tool/thin-bridge ablation arms. Full text read 19.09.2026 (INCOSE/Wiley, via FH Aachen library access); PDF at `bib/papers/dehart2024llm.pdf`.

- [7] DEHN, Simon ; SCHNÜRER, Simon ; JACOBS, Georg ; HÖPFNER, Gregor: Generating SysML v2 Models from Natural Language Requirements Using Large Language Models. In: *2025 IEEE International Symposium on Systems Engineering (ISSE)*, IEEE, Oktober 2025, p. 1–7. – Four-component structured prompting (role/goal description, ontology description, few-shot examples, natural-language requirements set) for NL→SysML v2 generation, evaluated on an automotive electric-window-opener use case with precision/recall/F1 for information extraction plus coverage, syntax pass/fail, and qualitative semantics scores across four incremental prompting setups (minimal to full structure), averaged over three runs each. Findings: LLMs extract system elements with high precision/good recall; the ontology component improves traceability but needs few-shot examples to reach syntactic correctness; few-shot examples are essential for syntactically valid SysML v2 output; structured prompting trades accuracy gains against higher token usage/runtime. Remaining weakness: unstable mapping of logical to physical elements. Primary reference point for the *Erstellen* use case and its element-level metrics; file-based generation without a modelling tool, unlike this thesis’s live-tool/MSoSA attachment point. PDF at [bib/papers/dehn2025nl2sysml.pdf](#).
- [8] QUALIS, Richard A.: Mitigating Hallucinations in SysML v2 Generation Using LLMs and a Tri-Layered Knowledge Graph Reasoning Framework. In: *2025 ACM/IEEE 28th International Conference on Model Driven Engineering Languages and Systems Companion (MODELS-C)*, IEEE, Oktober 2025, p. 357–366. – Tri-layer knowledge graph (Layer 1: reusable SysML pattern/annotation KG for nine diagram types; Layer 2: domain-specific KG for sectors like aerospace/automotive; Layer 3: auto-generated system-specific KG parsed from structured requirements) feeding reusable prompt templates to ground SysML v2 generation against hallucinated constructs; satellite-system case study evaluated on a manually curated ground-truth dataset (Average Structure Score, Average Manual Structure Score, Average Semantic Score, Consistency Ratio, prompt-type analysis). Produces structurally valid models reliably; semantic fidelity and output determinism remain limited. Motivates grounding the agent in the *live* model/-tool state rather than a hand-built static KG. Reference point for the *Erstellen* use case alongside [dehn2025nl2sysml](#), [wu2025sysforge](#), [bouamra2025systemp](#). PDF at [bib/papers/qualis2025hallucination.pdf](#).
- [9] QUAST, Vincent ; JACOBS, Georg ; DEHN, Simon ; HÖPFNER, Gregor: Enabling Humans and AI Systems to Retrieve Information from System Architectures in Model-Based Systems Engineering. In: *Systems* 14 (2026), Januar, No. 1, p. 83. <http://dx.doi.org/10.3390/systems14010083>. – DOI 10.3390/systems14010083. – GraphRAG multi-agent system (Supervisor Agent + Graph Query Agent) over a Neo4j knowledge graph parsed from SysML v2 textual notation by a custom Python parser and schema’d by the RFLP methodology; hybrid graph+vector indexing, Cypher

queries, query decomposition, following Peng et al.'s GraphRAG framework. Evaluated on a synthetic battery-EV reference model (4 abstraction levels) with a zero/one-hop vs. multi-hop QA dataset, comparing four LLMs (LLaMA-3.3-70B, Gemini-2.0-flash-lite, Gemini-2.0-flash, Gemini-2.5-flash-preview); best model reaches **93 % overall accuracy, 90 % on multi-hop questions**. Authors' own stated limitations: metamodel-specific parser, synthetic (not industrial-scale) model, no user studies, and explicitly "*lack of standardized benchmarks*" for QA-over-MBSE – an independent second data point for this thesis's benchmark-gap argument (alongside bouamra2025systemp). Primary reference for the *Abfragen* use case and its accuracy metric; the thesis queries the live tool/model state directly rather than a derived, pre-built graph. PDF at <bib/papers/quast2026graphrag.pdf>.

- [10] BOUAMRA, Yasmine ; YUN, Bruno ; POISSON, Alexandre ; ARMETTA, Frédéric: SysTemp: A Multi-Agent System for Template-Based Generation of SysML v2. In: *Advances in Practical Applications of Agents, Multi-Agent Systems, and Computational Social Science: The PAAMS Collection – 23rd International Conference, PAAMS 2025, Lille, France* Bd. 16031, Springer Nature Switzerland, Juni 2025 (Lecture Notes in Computer Science (LNAI)), p. 41–52. – Peer-reviewed version (supercedes arXiv:2506.21608). Multi-agent pipeline – SpecificationGeneratorAgent → TemplateGeneratorAgent → WriterAgent ↔ ParserAgent iterative repair loop – for NL-to-SysML-v2 textual generation. Ablation (GPT-4 Turbo, Claude 3.5 Sonnet, 5 toy bicycle-component scenarios): template scaffolding converges to zero syntax errors in 4/5 scenarios vs. 1/5 without it – syntax convergence only, no semantic evaluation. States that the only two available SysML v2 example corpora are the Pilot-Implementation test suite and GfSE/SysML-v2-Models, <150 scenarios total – corroborates this thesis's Apollo 11 (primary) / GfSE (secondary) split. Conclusion, quoted directly: "assessing the quality of generated SysML v2 models remains constrained by the absence of standard benchmarks" – the gap citation for this thesis's benchmark contribution. PDF at <bib/papers/bouamra2025systemp.pdf>.
- [11] AL-SHAMI, Haitham ; MALIK, Rohail ; ALA-LAURINHO, Riku ; VEPSÄLÄINEN, Jari ; VIITALA, Raine: Automated Semantic Fault Localization in SysML v2: A Human-in-the-Loop Framework Using Knowledge-Graph Augmented LLMs. In: *INCOSE International Symposium* 36 (2026), Juli, No. 1, p. 1274–1286. <http://dx.doi.org/10.1002/iis2.70189>. – DOI 10.1002/iis2.70189. – Fine-tuned small language model plus a domain KG of physical compatibility rules that localises and repairs *semantic* faults which pass SysML v2 syntax validation; synthetic faults injected via the KG, output as unified-diff patches, 1 184 test samples, repair rate <3 % → >91 %. Direct methodological precedent for the fault-injection design of the *Korrektur* use case and for the syntactic/semantic fault distinction. PDF at <bib/papers/alshami2026faultloc.pdf>.

- [12] BAZZAL, Mahmoud ; SCHLEIFER, Simon ; LUNGU, Adriana ; KRUSE, Benjamin ; GOETZ, Stefan ; KOLDEWEY, Christian ; TISSEN, Denis ; WARTZACK, Sandro J. ; DUMITRESCU, Roman: Model Context Protocol for MBSE: A Unified Interoperable Agentic Framework for Systems Modelling. In: *Proceedings of the Design Society* 6 (2026), Juli, p. 2721–2730. <http://dx.doi.org/10.1017/pds.2026.10630>. – DOI 10.1017/pds.2026.10630. – [TOP 3] Fraunhofer IEM / Paderborn HNI / FAU / Audi framework implementing MCP servers directly on the standard SysML v2 API (C#.NET, official reference API + database in Docker; open-source at github.com/hni-ase/SysMLV2-mcp). Three tool classes – creation, modification, analysis – mirror this thesis’s Create/Patch/Query use cases. Demonstrated, not benchmarked: two qualitative DSRM case studies (agentic use-case generation; signal-redundancy optimisation) adapted from prior workflow-based papers (Schleifer et al. 2025; Bazzal et al. 2025), no ground truth, no quantitative metrics – own discussion section names the “limitation on the number of MCP tools” and single-agent scope as open issues. Closest scientific competing work: delivers the *artifact class* (MCP-on-REST-API) the thesis builds on; the thesis supplies the missing empirical, ground-truth capability evaluation, and attaches inside the industrial tool (MSoSA) rather than the standard API alone. PDF at [bib/papers/bazzal2026mcpmbse.pdf](#).
- [13] DEMPSEY, Colin: *Native AI Integration for Model-Based Systems Engineering: Three Layers that Make It Work*. SEI Blog, Carnegie Mellon University, <https://www.sei.cmu.edu/blog/native-ai-integration-for-model-based-systems-engineering-three-layers-> September 2026. – [TOP 3] Controlled three-arm SysML v2 experiment, same 8 tasks at two scales (n=40 small, n=24 large), single model (Claude Opus 4.6): (1) baseline, one-shot, no tools; (2) CLI-only, generation + `syside check` validation loop; (3) full tooling, adds a repository-local MCP knowledge-retrieval server (curated Markdown, not an LLM or KG) plus 10 file-based “skills” (e.g. `model-sysml`, `fix-sysml`, `review-sysml`). Both iterative arms reach 0 syntax errors/warnings (baseline: 3.7/100 lines small, 25.6/100 lines large) – validation alone does not raise the project’s 29-measure “pattern” score, but full tooling does: 71.7 → 94.1 on large tasks. Worked example (4-drone survey swarm, 8 coordinated files, 97.2/100 pattern score, 0 validator errors) shows the score’s limit directly: separation and lost-link parameters, and several requested components, do not match the prompt, and analysis/verification elements are structurally present but not executable/evaluated – a clean validation run does not mean the model is correct. Entirely file-based (textual notation + Syside LSP); their MCP server is a knowledge-retrieval service, not a model repository or live tool interface – the thesis’s point of differentiation (live model inside an industrial tool, MSoSA). Closest *evaluated* work and the direct template for this thesis’s own

three-arm ablation (no tool / thin bridge / semantic bridge). PDF (archived readable-text capture, no native PDF export exists) at `bib/papers/sei-native-ai.pdf`.

- [14] ISO/IEC/IEEE: Systems and software engineering – System life cycle processes / International Organization for Standardization. Version: Mai 2023. <https://www.iso.org/standard/81702.html>. 2023 (ISO/IEC/IEEE 15288:2023). – International Standard. – Defines Verification (conformance to architectural/technical requirements, incl. traceability and syntactic checks) and Validation (fulfilment of stakeholder needs in the operational environment) as distinct system life-cycle processes. The thesis's Verification/Validation use cases adopt these definitions directly. Paywalled standard – no PDF copy in `bib/papers/`, cite by standard number.
- [15] CIBRIÁN, Eduardo ; OLIVERT-INSERTE, Jose ; DíEZ-FENOY, Carlos ; MENDIETA, Roy ; LLORENS, Juan ; ÁLVAREZ-RODRÍGUEZ, Jose M.: Ensuring Semantic Consistency in SysML v2 Models Through Metamodel-Driven Validation. In: *IEEE Access* 13 (2025), p. 121444–121457. <http://dx.doi.org/10.1109/ACCESS.2025.3587786>. – DOI 10.1109/ACCESS.2025.3587786. – Despite the title, this is ISO-verification: metamodel-conformance checking (syntax, cardinality, cross-reference integrity) via a custom validator, not behavioral simulation or stakeholder-objective fulfilment. Paper explicitly contrasts its static, structural approach with simulation-based tools (SysON). Method: metamodel-driven rule checker + case studies across three domains. Audited 19.09.2026, verdict: Verification. PDF at `bib/papers/cibrian2025validation.pdf`.
- [16] MOLNÁR, Vince ; GRAICS, Bence ; VÖRÖS, András ; TONETTA, Stefano ; CRISTOFORETTI, Luca ; KIMBERLY, Greg ; DYER, Pamela ; GIAMMARCO, Kristin ; KOETHE, Manfred ; HESTER, John ; SMITH, Jamie ; GRIMM, Christoph: Towards the Formal Verification of SysML v2 Models. In: *Proceedings of the ACM/IEEE 27th International Conference on Model Driven Engineering Languages and Systems (MODELS Companion '24)*, 2024, p. 1086–1095. – Surveys five tool integrations (Gamma, MP-Firebird, Imandra, SAVVS, SysMD) applying model checking, theorem proving and SMT/contract-refinement to SysML v2 models – all proofs of conformance to formally stated properties/contracts, with explicit traceability framing. "Verification and validation" appears only as a stock keyword pair; no operational-scenario or stakeholder-objective content anywhere in the body. Audited 19.09.2026, verdict: Verification. PDF at `bib/papers/molnar2024formalverification.pdf`.
- [17] RATZKE, Axel ; KOCH, Johannes ; GRIMM, Christoph: Modeling and Analysis of System Models with Constraints in SysMLv2. In: *2025 20th Annual System of Systems Engineering Conference (SoSE)*, 2025, p. 1–6. – Compiles KerML/SysML v2 range-based constraints (`oneOf/anyOf/allOf`) into a Constraint

Satisfaction Problem and solves it via constraint propagation – explicitly static, explicitly excludes dynamic/behavioral analysis (“restrict constraint propagation to... the static structure... excluding dynamic behaviors”). Uses “valid/validation” only informally for constraint satisfaction. Audited 19.09.2026, verdict: Verification. PDF at [bib/papers/ratzke2025constraints.pdf](#).

- [18] LIPARULO, Leonardo ; PIERRI, Francesco: *Benchmarking AI Agents for Hardware Design Automation via MCP Tool Calling*. arXiv:2608.26199 [cs.AI]. – Methodological sibling in an adjacent domain: an MCP server that reproduces the state and dependency logic of a proprietary hardware-design tool, plus a benchmark of single-operation edits, multi-step dependency chains, invalid requests and misspelled prompts, evaluated on seven open-source models under ablations of system prompt, tool-description detail, context scope and single- vs. multi-agent architecture (metric: expected-call coverage). Findings directly informing the bridge/benchmark design: comprehensive tool descriptions consistently reduce failures; few-shot prompting can cause severe inaction in some models; cumulative context harms constrained models; multi-agent decomposition helps weak workers/long sessions at extra call cost. Its task taxonomy (single-op/chain/invalid/misspelled) maps onto the SysML v2 benchmark’s difficulty tiers. PDF at [bib/papers/liparulo2026hwmcp.pdf](#).
- [19] HASAN, Mohammed M. ; LI, Hao ; RAJBAHADUR, Gopi K. ; ADAMS, Bram ; HASSAN, Ahmed E.: *Model Context Protocol (MCP) Tool Descriptions Are Smelly! Towards Improving AI Agent Efficiency with Augmented MCP Tool Descriptions*. arXiv:2602.14878. – Empirical study of 856 tools across 103 MCP servers, using six description components from the literature to build a scoring rubric and formalise “smells” (operationalised by an FM-based scanner): 97.1 % of descriptions carry at least one smell, 56 % fail to state purpose. Augmenting all components raises task success by a median 5.85 pp and partial completion by 15 %, but increases execution steps by 67 % and regresses performance in 17 % of cases; compact component combinations often preserve reliability at lower token cost. Supplies a rubric for the bridge’s tool-description design and identifies a confound to control between the thin and semantic bridge arms: description quality must be held constant so the semantic arm does not win merely on better-written tools. PDF at [bib/papers/hasan2026mcpsmells.pdf](#).